

## Lab 4: Yelp Reviews

Kelly Xie

Apr 11, 2017

### 1. Summary Statistics


```
# a histogram of all restaurant ratings given by users

#install.packages("ggplot2")
library("ggplot2")

reviews = read.csv("~/Desktop/YelpReviews_20k.csv")
reviews = unique(reviews) #remove duplicate reviews

gg = ggplot(data=reviews, aes(reviews$stars)) +
  geom_histogram(binwidth = 1, aes(fill = ..count..)) +
  xlab("Restaurant Rating") +
  ylab("Frequency of Rating") +
  ggtitle("Frequency of Restaurant Ratings on Yelp")

gg
```



```
# calculate the number of reviews the average restaurant in this sample received

#install.packages("dplyr")
library("dplyr")

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

byrestaurant = reviews %>%
  group_by(business_id) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
summarise(byrestaurant, "Average Number of Reviews by Restaurant" = mean(count))
```

```

## # A tibble: 1 × 1
##   `Average Number of Reviews by Restaurant`
##                                     <dbl>
## 1                                     2.420441

# calculate the number of reviews the average user has contributed in this sample

byuser = reviews %>%
  group_by(user_id) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
summarise(byuser, "Average Number of Reviews by User" = mean(count))

## # A tibble: 1 × 1
##   `Average Number of Reviews by User`
##                                     <dbl>
## 1                                     1.152578

# on average, do GoodForLunch restaurants receive a greater number of reviews

goodforlunch = reviews %>%
  subset(GoodforLunch == "True") %>%
  group_by(business_id) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
summarise(goodforlunch, "Average Number of Reviews (GoodForLunch)" = mean(count))

## # A tibble: 1 × 1
##   `Average Number of Reviews (GoodForLunch)`
##                                     <dbl>
## 1                                     2.189089

notgood = reviews %>%
  subset(GoodforLunch == "False") %>%
  group_by(business_id) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
summarise(notgood, "Average Number of Reviews (NOT GoodForLunch)" = mean(count))

## # A tibble: 1 × 1
##   `Average Number of Reviews (NOT GoodForLunch)`
##                                     <dbl>
## 1                                     2.590708

```

***On average, restaurants marked Good For Lunch receive 0.4 (15.5%) fewer reviews than restaurants that are not.***

```
# on average, do GoodForLunch restaurants receive a higher number of stars
```

```
GFL = reviews %>%  
  subset(GoodforLunch == "True") %>%  
  group_by(business_id) %>%  
  summarise(rating = mean(stars)) %>%  
  arrange(desc(rating))  
summarise(GFL, "Average Rating (GoodForLunch)" = mean(rating))  
  
## # A tibble: 1 × 1  
##   `Average Rating (GoodForLunch)`  
##                               <dbl>  
## 1                               3.617724  
  
notGFL = reviews %>%  
  subset(GoodforLunch == "False") %>%  
  group_by(business_id) %>%  
  summarise(rating = mean(stars)) %>%  
  arrange(desc(rating))  
summarise(notGFL, "Average Rating (NOT GoodForLunch)" = mean(rating))  
  
## # A tibble: 1 × 1  
##   `Average Rating (NOT GoodForLunch)`  
##                               <dbl>  
## 1                               3.605655
```

***On average, restaurants marked Good For Lunch on Yelp are rated 0.012 (0.33%) higher than restaurants that are not.***

## 2. Exploratory Text Analysis

```
#convert reviews to text corpus  
#install.packages("tm")  
library("tm")  
  
## Loading required package: NLP  
  
##  
## Attaching package: 'NLP'  
  
## The following object is masked from 'package:ggplot2':  
##  
##   annotate  
  
corp.original = VCorpus(VectorSource(reviews$text))  
  
#clean and reprocess the text  
corp = tm_map(corp.original, removePunctuation)  
corp = tm_map(corp, removeNumbers)  
corp = tm_map(corp, content_transformer(removeWords), stopwords("SMART"), lazy
```

```

=TRUE)
corp = tm_map(corp, content_transformer(tolower), lazy=TRUE)
corp = tm_map(corp, content_transformer(stemDocument), lazy=TRUE)
corp = tm_map(corp, stripWhitespace)

#generate a document-term matrix
dtm = DocumentTermMatrix(corp)
m = as.matrix(dtm)

#get fifteen most frequently appearing words among all reviews
word.freq = colSums(m)
word.freq = sort(word.freq, decreasing=TRUE)
as.data.frame(word.freq[1:15])

##           word.freq[1:15]
## food                15778
## good                13901
## place              13769
## order              10041
## great              9642
## veri               8774
## time               8479
## servic             8032
## tri                6432
## back               5876
## realli             5870
## restaur            5642
## friend             5311
## love               5199
## onli               4388

#generate a word cloud using the document-term matrix (max 100 words)
#size of the word correlates to its frequency in the review

#install.packages("wordcloud")
library("wordcloud")

## Loading required package: RColorBrewer

wordcloud(names(word.freq), word.freq, scale = c(4, .5),
          max.words = 100, colors = brewer.pal(6, "Dark2"), random.order = FALSE)
LSE)

```



### 3. Text Analytics and Prediction

```
#get number of unique terms in document-term matrix
dim(dtm)
```

```
## [1] 19988 34255
```

***There are 34255 unique terms in the document-term matrix.***

```
#narrow the list down to the 200 words with the most predictive power
```

```
dtms = removeSparseTerms(dtm, .990) #remove sparse terms with .990 threshold
dtm.matrix = as.matrix(dtms)
```

```
#calculate correlation matrix between document-term matrix and goodforlunch
corr = cor(as.numeric(as.logical(reviews$GoodforLunch)), dtm_matrix)
absCorr = abs(corr) #get absolute value of correlations
```

```
#keep 200 terms with highest correlation magnitudes (both pos and neg)
top200 = order(absCorr, decreasing=TRUE)[1:200]
top200words = colnames(absCorr)[top200]
```

```

#create new document-term matrix with these terms
newDTM.df = as.data.frame(cbind(GoodForLunch = as.numeric(as.logical(reviews$
GoodforLunch)),
                                dtm_matrix[,top200words]))
newDTM.m = as.matrix(newDTM.df)
dim(newDTM.m)

## [1] 19988    201

```

***There are 201 unique terms in the new document-term matrix.***

*#generate a wordcloud where the size corresponds to the correlation strength  
#of the top 20 positive and negative words*

```

#get top 20 positive words
top20pos = order(corr, decreasing=TRUE)[1:20]
top20poswords = colnames(corr)[top20pos]
pos.df = as.data.frame(cbind(term = top20poswords, corr = corr[top20pos]))

#get top 20 negative words
top20neg = order(corr)[1:20]
top20negwords = colnames(corr)[top20neg]
neg.df = as.data.frame(cbind(term = top20negwords, corr = corr[top20neg]))

#form wordcloud
wordcloud(words = c(as.character(pos.df$term), as.character(neg.df$term)),
          freq = c(as.numeric(as.character(pos.df$corr)), abs(as.numeric(as.c
haracter(neg.df$corr)))),
          scale = c(2.5, .5),
          colors = c(rep("green",20), rep("blue",20)),
          ordered.colors = TRUE, random.order = FALSE, random.color = FALSE)

```



**Legend: Positive term = green; Negative term = blue**

*#Partition the matrix into training and test rows so you can use the test data to evaluate your model performance. Set the last 20% of your rows aside for testing, and use the first 80% to build your model as specified below*

```
traindata = newDTM.df[1:(.8*nrow(newDTM.df)),]
testdata = newDTM.df[-(1:(.8*nrow(newDTM.df))),]
```

*#Fit a logistic regression model to the selected variables in the training data.*

```
model = glm(GoodForLunch ~ ., data = traindata, family = binomial)
model
```

```
##
```

```
## Call: glm(formula = GoodForLunch ~ ., family = binomial, data = traindata)
```

```
##
```

```
## Coefficients:
```

## (Intercept)	reserv	dessert	breakfast	sandwich
## -3.180e-01	-1.040e+00	-2.614e-01	-4.745e-01	3.041e-01
## wine	night	servic	lunch	burger
## -2.789e-01	-2.090e-01	-1.481e-01	2.298e-01	2.034e-01
## chicken	steak	filet	fri	dinner
## 1.608e-01	-3.314e-01	-1.150e+00	1.441e-01	-9.434e-02



##	waiter	server	noodl	coffe	dog
##	-1.638e-01	-9.503e-02	4.819e-01	-3.040e-01	4.054e-01
##	view	brunch	chines	pancak	atmospher
##	-1.034e+00	-4.643e-01	2.770e-01	-6.119e-01	-2.571e-01
##	taco	drink	beauti	experi	chocol
##	1.625e-01	-8.167e-02	-3.054e-01	-1.262e-01	-2.507e-01
##	rice	veggi	butter	birthday	attent
##	1.679e-01	3.729e-01	-3.812e-01	-1.118e-01	-1.260e-01
##	cocktail	hot	ambianc	authent	pho
##	-2.137e-01	1.010e-01	-4.490e-01	2.608e-01	3.045e-01
##	fast	potato	celebr	tabl	vega
##	2.753e-01	-1.723e-01	-6.848e-01	7.062e-02	-1.339e-01
##	clean	cake	entre	chef	egg
##	3.272e-01	-1.498e-01	-2.110e-01	-1.476e-01	-3.018e-01
##	cours	mexican	dine	salmon	start
##	-8.707e-02	1.727e-01	1.217e-02	-3.050e-01	-9.304e-02
##	bbq	happi	glass	morn	great
##	1.282e-01	-1.242e-01	3.953e-02	-4.182e-01	-3.540e-02
##	bar	seat	restaur	bun	light
##	-1.284e-02	2.573e-02	-5.525e-02	1.837e-01	-1.752e-01
##	appet	date	scallop	today	perfect
##	2.945e-02	-2.970e-01	-2.274e-01	3.045e-01	2.741e-02
##	combo	cream	beef	french	bacon
##	6.161e-02	-1.031e-01	9.777e-02	-2.069e-01	-3.221e-01
##	bartend	amaz	crab	pull	excel
##	-1.358e-01	-9.025e-02	-1.370e-01	2.012e-01	-1.254e-01
##	group	crave	room	toast	burrito
##	-2.680e-02	3.739e-01	-5.205e-02	-1.337e-01	2.351e-01
##	show	bowl	parti	end	sat
##	-3.007e-01	1.001e-01	-8.010e-03	-1.881e-02	-9.502e-02
##	bottl	counter	meat	place	bean
##	-1.417e-01	1.624e-01	1.134e-01	8.718e-02	3.758e-02
##	calamari	prime	music	mash	dish
##	-3.855e-01	-2.237e-01	-2.186e-01	-1.721e-01	-9.010e-02
##	pork	thai	veri	hour	main
##	-2.512e-02	1.487e-01	-2.286e-02	1.491e-02	3.460e-02
##	salsa	sushi	turkey	diner	menu
##	9.205e-02	-1.084e-01	2.421e-01	-2.216e-01	4.576e-03
##	present	fantast	hous	nice	finish
##	-2.610e-02	-1.588e-01	-6.437e-02	-2.461e-02	-1.462e-01
##	soda	sauc	onion	waffl	deliveri
##	2.572e-01	9.441e-02	8.538e-02	-1.930e-01	2.112e-01
##	charlott	share	meal	lettuc	tortilla
##	5.867e-01	-7.131e-02	1.188e-02	-1.448e-02	2.785e-01
##	felt	tuna	decor	wait	quick
##	-9.655e-02	-1.269e-01	1.199e-02	-4.023e-03	2.092e-01
##	arriv	seafood	bread	guy	pickl
##	2.179e-01	-1.957e-01	-7.388e-02	1.684e-01	9.847e-02
##	healthi	set	app	hotel	cut
##	3.402e-01	-2.904e-02	-4.910e-01	-3.778e-02	-1.827e-01

```

##      patio      impress      list      knowledg      fruit
## -7.831e-02 -1.430e-01 -4.464e-02 4.713e-02 5.609e-03
##      curri      lamb      expens      ring      worth
## -3.642e-02 -3.121e-01 -2.834e-01 1.076e-01 -1.515e-01
##      add      cool      strip      dark      late
## 1.464e-01 -2.404e-01 -5.616e-02 -1.591e-01 -2.213e-01
##      oyster      joint      pasta      short      pour
## -8.602e-02 1.448e-01 -2.594e-02 1.307e-02 -1.976e-01
##      hostess      box      weekend      split      fresh
## 3.579e-02 3.393e-01 -4.483e-02 -2.149e-01 1.692e-01
##      lobster      comfort      relax      casino      locat
## 1.504e-01 -1.660e-01 -2.586e-01 -3.807e-02 4.347e-02
##      run      guest      vibe      leg      oliv
## 2.171e-01 -9.527e-02 -2.371e-01 4.525e-02 -1.924e-01
##      banana      saturday      spinach      water      soup
## 1.253e-01 -5.698e-03 -1.179e-01 -3.324e-02 1.153e-02
##      outstand      enjoy      friday      befor      waitress
## 1.183e-02 8.100e-02 3.038e-02 1.911e-03 -5.483e-02
##      fun      desert      outdoor      greasi      fine
## -1.239e-01 -1.917e-01 -2.245e-01 1.974e-02 -1.141e-01
##      floor      plate      surpris      rich      book
## -2.802e-01 -2.806e-02 -1.066e-01 1.531e-01 -1.647e-01
##      overal      delight      valley      talk      eat
## -9.203e-05 -1.764e-01 5.799e-01 -7.310e-02 1.021e-01
##      trip
## -1.549e-01
##
## Degrees of Freedom: 15989 Total (i.e. Null); 15789 Residual
## Null Deviance: 21300
## Residual Deviance: 18650 AIC: 19050

#A positive coefficient positively predicts that a restaurant is good for Lunch.
#A negative coefficient suggests a restaurant would not be good for lunch.

#Use the coef command to access top positive and negative words from the model.
coef = coef(model)[-1]
pos.terms = coef[coef>0]
top.pos = sort(pos.terms,decreasing=T)[1:15]
top.pos

## charlott valley noodl dog crave veggi healthi
## 0.5866726 0.5798723 0.4819046 0.4053690 0.3738928 0.3728794 0.3401948
## box clean today pho sandwich tortilla chines
## 0.3393160 0.3271950 0.3045467 0.3045441 0.3041261 0.2784771 0.2770426
## fast
## 0.2752966

```

```

neg.terms = coef[coef<0]
top.neg = sort(neg.terms)[1:15]
top.neg

##      filet      reserv      view      celebr      pancak      app
## -1.1499466 -1.0404650 -1.0341077 -0.6848086 -0.6118874 -0.4909538
## breakfast brunch ambianc morn calamari butter
## -0.4745300 -0.4643414 -0.4489867 -0.4181817 -0.3855215 -0.3812082
##      steak      bacon      lamb
## -0.3313668 -0.3221268 -0.3120732

#Produce a word cloud that separates the top 15 positive words and top 15 neg
ative words.

poswords = tibble::rownames_to_column(as.data.frame(top.pos), var="term")
negwords = tibble::rownames_to_column(as.data.frame(top.neg), var="term")

#form wordcloud
wordcloud(words = c(poswords$term, negwords$term),
          freq = c(poswords$top.pos, abs(negwords$top.neg)),
          scale = c(4.5, .5),
          colors = c(rep("red",15), rep("purple",15)),
          ordered.colors = TRUE, random.order = FALSE, random.color = FALSE)

```



**Legend: Positive words = red; Negative words = purple**

*#Using the model you have generated, choose a probability threshold to maximize accuracy and classify the restaurants in your training data as 1 or 0 according to whether they are GoodForLunch.*

```
traindata$predict_val = predict(model, type="response")
traindata$gfl_predicted = traindata$predict_val > 0.5
```

*#How well does this model perform on the training data in terms of classification accuracy (i.e. the percentage of GoodForLunch values that you get correct)?*

```
accuracy = sum(traindata$gfl_predicted == traindata$GoodForLunch) / nrow(traindata)
accuracy
```

```
## [1] 0.6843027
```

***With a probability threshold of 0.5, this model has a 68.43% classification accuracy for predicting GoodForLunch in the training data.***

*#Predict values for GoodforLunch in your test data.*

```
testdata$predict_val = predict(model, newdata = testdata, type = "response")
testdata$gfl_predicted = testdata$predict_val > 0.5
```

*#How well does the model perform in terms of classification accuracy (i.e. the percentage of GoodForLunch values that you get correct)?*

```
accuracy.test = sum(testdata$gfl_predicted == testdata$GoodForLunch) / nrow(testdata)
accuracy.test
```

```
## [1] 0.6783392
```

***With a probability threshold of 0.5, this model has a 67.83% classification accuracy for predicting GoodForLunch in the test data.***